

OutLoud

Practise speaking English. Nobody is listening but you.

Product Management Capstone · Ankita Choudhary · 18 September 2026

[instant-capture-show.lovable.app](#)

WHERE THIS STARTED

**He wrote the whole script on his cue card.
Then skipped the words that felt too hard.**

“The others were laughing because he skipped the difficult word.”

The teacher saw the script

His assignment marks were cut

The class laughed

At the one word he skipped

He withdrew

Decided not to present again at all

Fear → concealment → exposure → withdrawal. That last step is what makes this a product problem.

Someone who already knows the words

Who this is for

- College students with graded spoken assessments — presentations, vivas, group discussions
- Reads English job postings and writes emails in English without difficulty
- Freezes when there are people in the room
- Has a spoken event within the next six weeks
- Smartphone only. No laptop, no fixed desk.

Who this is not for

- Absolute beginners — they need vocabulary, this user needs repetitions
- Anyone wanting accent training or certification
- People who want a human tutor — that is the thing being avoided
- "Everyone who isn't confident in English" — that is not a segment, it is the absence of one

Six conversations, five in segment

- 1 Kartik — B.Com student
- 2 Anuradha — BMF student
- 3 Dhruv — engineering
- 4 Shubham — engineering
- 5 Kapil — working professional
- 6 Puneet — contrast case, screened out

Recruited across fields deliberately

I did not want to assume the segment was defined by the kind of college someone attends. Screened on behaviour instead: can read and write English, has a recent incident, not a beginner.

Stated honestly

The brief asks for 8–10 conversations. I completed 6 — five in segment, plus one contrast case who was screened out.

Every frequency in this deck is out of the five in-segment participants. The sample skews to current students, and there is no ITI or polytechnic representation.

What five people actually said

FINDING	COUNT	WHY IT MATTERS
Last English-speaking moment was a class or group presentation	5 / 5	Not a job interview — the pain arrives years earlier, and on a timetable
Read from cue cards during it	most	Concealment is the default coping strategy
Will not practise in front of a human / do not want to be judged	3 / 5	The core finding
Comfort speaking English to a stranger tomorrow	mean 1.6	Out of 5. Three at the floor, none above 3
Abandoned a paid human-tutor option	≥1 named	PlanetSpark — per-session pricing, "very expensive"
Withdrew from a future speaking opportunity	1 / 5	The consequence that makes this worth solving
Has a dated speaking event in the next six weeks	1 / 5	October internship interview — the retention hook

THE DAY-0 BASELINE

“How comfortable would you feel speaking English to a stranger tomorrow?”


1.6

mean, out of 5

3 of 5

sat at the absolute floor

Nobody above a 3.



A different question from the interview, and it lands on the same finding: the one person willing to try qualified it on audience size — “but not in front of many people.” The variable is not the language. It is who is watching.

STEP 1

Three insights

Each one passes the same test: if the opposite were true, v1 would look different.

The fear is of the witness, not the language.

THE EVIDENCE

- 3 of 5 will not practise in front of a human
- The classmates laughed at the skipped word
- Comfort with a stranger averages 1.6/5 — and the only person above a 2 said “but not in front of many people”
- Nobody in the sample asked for better teaching
- And a sixth person, who had already solved it, built this himself — practising English on calls to Amazon support

SO V1 DOES THIS

v1 removes every observer: no human, no score, no grade, no leaderboard, no login. And a retake leaves no record of the attempt before it — failing has to be untraceable.

Concealment is the coping strategy — and it is what breaks.

THE EVIDENCE

- He wrote the full script on the cue card
- Then skipped the hard words anyway
- The script is exactly what got him caught, and what got him laughed at
- Most of the sample read from cue cards

SO V1 DOES THIS

v1 makes the script impossible: a prompt card, 15 seconds to think with an explicit instruction not to write it down, then 60 seconds of unscripted speech. The product rehearses the one thing they currently avoid.

The alternatives fail on economics and invisible progress — not teaching.

THE EVIDENCE

- A human in the loop costs ₹200–500 an hour
- Which forces long booked sessions — colliding with a student who wants five minutes on a bus
- “No definite curriculum” — a direct quote. They paid, practised, could not see what improved, and left

SO V1 DOES THIS

v1 is five minutes, free, phone-only, and keeps a retrievable history of what was practised. AI is not merely the cheap option — it is the only thing that makes a five-minute session economically possible.

THE PERSONA

Rohan

Built from interview 2. Second-year student.

SITUATION

Reads English job postings without difficulty, writes his assignments in English, follows lectures in English. Has never held a two-minute spoken conversation in English without preparing the words first.

THE RECENT EVENT

Wrote his full presentation script onto a cue card. Skipped the words that felt too hard. The class laughed. The teacher spotted the script and cut his marks.

GOAL

Get through the next presentation without being laughed at — and sit an internship interview without freezing.

BIGGEST FRUSTRATION

There is nowhere to be bad at this. Coaching means a room of peers. A paid tutor is still a person watching, at ₹200–500 an hour. His own home has people in it.

WHAT HE ASKED FOR

Five minutes on his phone, whenever he wants, including on the bus. Nobody listening. Something that remembers what he did last time.

WHAT WOULD MAKE HIM QUIT

A score. Anything that feels like being graded reintroduces exactly what he is running from.

SECONDARY RESEARCH

Every alternative has a person in it, or teaches the wrong thing

OPTION	PRICE	SPEAKS OUT LOUD?	A HUMAN PRESENT?	WHY IT FAILS THIS USER
PlanetSpark	per session, "very expensive"	Yes	Yes — 1:1 tutor	The observer is the product. Per-session pricing forces long booked slots.
Local coaching centres	₹2,000–5,000 / month	Sometimes	Yes — teacher + peers	A room of peers is the exact situation being avoided.
Private tutors	₹200–500 / hour	Yes	Yes	Cost compounds monthly; still a person watching you fail.
SpeakX	₹299 / month	Yes, to AI	No	Closest competitor. ~30% monthly paid retention.
Supernova AI	₹499 / month	Yes, to AI	No	Same category. Auto-renewal and deletion complaints.
MySivi / PeerUp	₹1–499	Yes, to a peer	Yes — a co-learner	Peer matching reintroduces a listener, and fails when nobody connects.
EngVarta	₹108 / session	Yes	Yes — live expert	Best problem framing of any incumbent, but a human caps every session. No free tier.
Duolingo	freemium	Barely — mostly tapping	No	Teaches A1 vocabulary to beginners. This user knows the words.
Free YouTube content	free	No	No	Watching someone speak English does not build the ability to speak it.

Two corrections: Hello English was removed from the Play Store in Feb 2025 and Enguru is effectively defunct — both appear in most analyses of this space and neither is live.

THE MARKET'S OWN PROOF

70%

of SpeakX's paying subscribers
churn every month

10M+

Play Store installs

₹45 Cr

FY26 revenue

\$20M+

raised

1 day

CAC payback

MS Dhoni

strategic investor, 1 Sep 2026

India's largest AI English-speaking app cannot get people to keep practising — and with a one-day CAC payback, its economics do not require anyone to learn. Retention is the unsolved problem. Retention is exactly what this project is measured on.

STEP 2 • THE PROBLEM

College students in India who read and write English fluently need to speak it out loud under pressure — in presentations, vivas and interviews — but every available place to practise has a person in it who might judge them.

So they conceal instead of practising: they script what they will say, skip the words that feel hard, get caught, and withdraw from the next opportunity.

We will solve this by giving them a five-minute, unwitnessed practice loop on their phone, where speaking badly leaves no trace.

A COMPLETED PRACTICE SESSION

≥20 seconds of unscripted speech in response to a prompt, plus viewing the coaching response. Does not count: opening the app · reading aloud from text · tapping through without speaking.

One five-minute loop, five steps

1

The card

One prompt, matched to what they said is coming up. “Describe one problem your project solves, in simple words.”

2

15 seconds to think

A countdown, and the instruction: “Think about what you'll say. Don't write it down.”

3

Speak, 60 seconds

One large button, a waveform, a countdown. No live transcript, no corrections, no score climbing.

4

Three things back

One thing you did well • one thing to try next time • “You said... / You could say...”

5

Try it again


Unlimited retakes — and the earlier attempt is stored nowhere. This button is the product.

Total: about five minutes — which is what the participants asked for, in those words.

THE THREE SCREENS THAT DECIDE WHETHER THEY COME BACK


Built and live on a mobile browser

Two real phone screens from users on 18 September 2026, plus the Speak screen as a wireframe.




Speaking

A waveform and a countdown — and deliberately nothing that judges them mid-sentence.



The response

One strength, one thing to try, one rewrite. No score anywhere.



Return and history

A real user's phone. His class presentation is in 18 days, and the session he did is saved below it.

What OutLoud deliberately will not assess

Accent

There is no defect to correct. Accent-shaming is the mechanism of the original harm.

Grammar, in v1

Their written grammar is fine. Real-time production is the gap.

Pronunciation score

Requires a numeric judgment, which reintroduces the grade.

Overall English level

Any level label becomes an identity a nervous user carries around.

Comparison with others

The whole product is the absence of an audience.

Whether the content was right

That is their teacher's job. OutLoud coaches the speaking, not the subject.

Exactly one strength, one suggestion, one rewrite. Never more.

My participants are not short of correction. They have a teacher who cut their marks and a class that laughed. A fourth critic on their phone would reproduce, in private, the exact experience they are trying to escape — and the one participant who explained why AI appealed to him said it was because nobody is judging. If OutLoud judges, it has no reason to exist.

Hard rules in the system prompt

- No score, grade or percentage — ever
- Never the words wrong, mistake, error, poor, incorrect
- Never a comment on accent, never grammar terminology
- Transcription errors are the system's fault, not the user's
- Strength always first. Under 120 words.

THE TRADE-OFF

OutLoud corrects less English per session than every competitor.

I am optimising for the user opening it on day 25, not for maximum improvement on day 1.

If Week 4 completion comes in strong and self-reported willingness has not moved, that trade-off was wrong.

What I measured, and what I could not

Week 4 practice completion is not measurable in this window.

OutLoud shipped and acquired its first users on 17–18 September 2026. Four calendar weeks have not elapsed. Rather than present a modelled number as a result, this is stated as the single biggest limitation in the evidence base.

METRIC	TARGET	ACTUAL	STATUS
Day-0 comfort, mean of 5	—	1.6 / 5	Measured
Link opens	—	[fill in]	Measured
First-session completion	60%	[fill in]	Measured • the strongest predictor of Week 4
Users with ≥1 completed session	50	[fill in]	One day of acquisition
Post-use comfort, same question	2.6	[fill in]	Measured
Week 1 completion	—	not yet elapsed	Unmeasurable
Week 4 completion	40%	not measurable	The evidence gap

Drop-off hypothesis, written before any data arrived: the largest drop comes in Week 2 or 3, immediately after the user's graded event passes — because Week 1 opens are driven by a deadline, and when the deadline is gone so is the reason.

How the first users responded

7

People approached

5

Opened the link

2+

Completed a session

1

Came back for a 2nd

71% opened a cold link with no download and no signup. Two users completed a session on their own phone and one did a second, unprompted, five minutes later. These are floors, not totals — v1 shipped without analytics, so anyone who used it and didn't send a screenshot is invisible here. That is a real mistake and the first thing I would fix.

What came back

5 / 5

app experience, from every tester

Worth reporting, not worth much. It measures satisfaction, not willingness — and every tester knew me and knew I built it. The honest reading: nobody disliked it and nobody got stuck, which is a real signal about first-session friction and nothing more. The willingness re-ask was not run. That is a gap, not a result.

The 20-second pitch I used

"You know when you have to present in class and you already know the answer, but the words don't come out? I made something for that. No download, no signup. Nobody hears it but you — and if it goes badly you press try again and there's no record of it."

What could kill this

RISK	LIKELIHOOD	MITIGATION
The AI discourages the exact user it is for	High	Hard prompt constraints: one suggestion max, banned words, no accent comment, strength first
Speech recognition mishears an Indian accent and "corrects" something said correctly	High	No pronunciation or accent scoring at all. Transcription artefacts treated as the system's error
Browser mic permission fails on mid-range Android	High	Tested on a real device. Permission asked only on first tap, with a reason. Typed fallback shipped
Retention collapses when the graded event passes	High	History and countdown are in v1 for this. The manufactured-deadline experiment is designed and ready
Nobody speaks out loud at all — they see the mic and leave	Medium	Session starts immediately; 15-second timer lowers the stakes; silence handled as normal
Research base too thin — n=5, no ITI or polytechnic	High	Stated openly, nothing extrapolated. Next round targets the missing segments

Trade-offs I chose: learning speed, a measurable progress score, absolute beginners as a market, and the entire login-based retention toolkit — all given up to keep the product anonymous and short.

THE DECISION I AM LEAST CONFIDENT ABOUT

Showing no score at all.

The reasoning traces to the research: a score is a verdict, my participants are running from verdicts, and 3 of 5 said they would not practise in front of anyone judging them.

But there is an argument on the other side I cannot dismiss. Their complaint about the paid apps was “no definite curriculum” — no visible progress. A score is the crudest and most legible form of visible progress there is, and I removed it. I am asking a saved history list to do a job a number does better.

So the honest position is that I may have over-corrected: I optimised entirely against one fear, using a second fear as justification, when the two may trade off against each other rather than point the same way.

The v2 plan reflects the uncertainty rather than resolving it — an opt-in score, unlocked only after session three, once the user has evidence this place is safe. I do not yet know whether that is the right answer or a compromise that satisfies nobody.

Next

What surprised me

- The speaking moment was a class presentation, not a job interview. The pain arrives years earlier than the brief assumed — and it recurs on a timetable.
- Nobody asked for better teaching. They asked for privacy.
- The cue-card script — the coping strategy — was the thing that got him caught.
- Someone had already built OutLoud by hand. A sixth participant taught himself English by calling Amazon support to talk to a stranger who could not see him, then asking them to rate his English out of 10. He did not want a teacher. He wanted a listener who was not watching.

The experiment I would run first

Manufacture the next deadline. On session two, ask the user to name their next speaking situation and its date, and count down to it. Half the cohort gets the prompt, half does not.

If the countdown group retains better, retention is a scheduling problem and the roadmap chases calendars, not gamification. If it makes no difference, the drop-off is in the loop itself and the feedback design needs rework. Either result kills a direction.

THE ROADMAP

Immediately

Run the cohort for four real weeks and measure the North Star properly

Next

Widen research to ITI and polytechnic students — the segment missing from n=5

Then

Opt-in score after session three, and prompt packs for specific vivas

Later

Hindi and regional-language scaffolding in the interface, and a 3-minute full-presentation mode

**“With AI I can be chill.
I know nobody is judging me.”**

Kartik, interview 1

OutLoud gives him somewhere to be bad at it, where nobody is watching.

instant-capture-show.lovable.app · Ankita Choudhary · 18 September 2026