

Airbnb India: Growth Loops for a Low-Take-Rate Market

Product-led growth initiatives, product-embedded loops and flywheels, MVP designs, and a revenue model built for India.

CORE THESIS

Airbnb India's binding constraint is **low trip frequency on trust-gated demand, sitting on thin and unprofessional supply, in a market that will not pay a 15.5% take rate on a ₹4,000 room night**. The answer is not to raise the take rate. It is to make the product manufacture trust, do the work the org currently does by hand, and monetise **the trip** rather than the room-night.

CONTENTS

What this document contains

1 · Executive summary	3
2 · Context & diagnosis — the four constraints	5
3 · Product audit — what already exists vs. what is missing	7
4 · PLG initiatives for cross-functional growth	8
Brief requirement (i)	
5 · Product-embedded growth loops & the India flywheel	11
Brief requirement (ii)	
6 · MVP designs — hypotheses, prototypes, success metrics	15
Brief requirement (iii)	
7 · Revenue model for India	23
Brief requirement (iv)	
8 · Prioritisation & roadmap	26
9 · Risks, assumptions & limitations	28
10 · Frameworks, research method & references	30
Brief requirement (v)	

A NOTE ON METHOD

Before proposing anything, I audited what Airbnb already ships. Two of the ideas I initially wanted to propose — reciprocal host–guest reviews and guest-curated collections — turned out to exist already. Rather than drop them, I repositioned them around the part that is genuinely broken: the reciprocal review exists but has *no consequence attached*, and curation exists but has *no creator economics*. Section 3 shows that audit in full. I would rather be marked on a correct diagnosis of an existing feature than on a confident proposal of one that shipped years ago.

COMPANION PROTOTYPE

All four MVPs in Section 6 are implemented as a single working, clickable prototype (StayCircle — a concept build, not affiliated with Airbnb, Inc.). Every screenshot in this document is taken from it. It includes a demo-control panel that toggles each feature on and off, which is how the control-vs-variant image pairs in Section 6 were produced.

SECTION 1

Executive summary

Airbnb wants to be an end-to-end travel platform built on human connection rather than commodity rooms. In India, the economics of that ambition are hostile: bookings are low-value, infrequent, trust-constrained, and contested by incumbents who charge more and offer deeper inventory. This document argues that the way through is not a better funnel but a better *loop* — and not a higher take rate but a different one.

The diagnosis

Four constraints bind Airbnb India, and every recommendation here attacks at least one:

Trust

LISTING ≠ REALITY
IS THE TOP GUEST
OBJECTION

Frequency

1–2 TRIPS PER
USER PER YEAR

Supply

THIN,
UNPROFESSIONAL,
LEAKAGE-PRONE

Take rate

15.5% ON A LOW-
VALUE BOOKING

The five moves

MOVE	WHAT IT IS	ATTACKS
Product-led ops	Convert the manual work Supply Ops, Support, Marketing and Trust & Safety do today into product surface and shared signals — an AI listing copilot with a signal-ranked ops queue, WhatsApp-native multilingual self-serve support, programmatically generated trust pages, and a verified-identity layer.	Cost structure; supply
Proof of Stay	Let guests attach verified photos to reviews. Booking.com has had this since 2016 and Vrbo shipped it in October 2025; Airbnb has not. It is a live competitive gap and the highest-leverage trust fix available.	Trust
Trust Score	Reciprocal reviews already exist — but the guest's rating is invisible and carries no consequence, so it changes no behaviour. Attach real benefits (waived deposit, Instant Book) and it becomes a loop.	Trust; supply
Split & Invite	India travels in groups, yet a group booking creates exactly one account. Make the booking a multiplayer object settled over UPI and every co-traveller becomes an activated user at near-zero CAC.	Frequency; acquisition
Trip, not room	A three-layer revenue model: a take rate that <i>falls</i> as a host earns quality and loyalty, ancillary trip services that monetise incremental spend, and subscriptions on both sides.	Take rate

The counter-intuitive claim

LOWER THE TAKE RATE AND EARN MORE

Section 7 models a booking where the headline take rate **falls from 15.5% to a blended 13%** — and contribution per booking still **risers roughly 10%**, because trip-services attach, subscription revenue and cheaper UPI rails more than replace the forgone commission, while Trust Score and AI support cut the cost to serve. That model excludes the volume effect from reduced direct-booking leakage, so it is the conservative case.

What I would build first

On a RICE basis (Section 8), **Proof of Stay** ranks first: it is cheap, it addresses the top guest objection, it is a verified gap versus competitors, and its output — genuine photographs of real Indian properties — is simultaneously a trust asset, free marketing creative and SEO inventory. One feature, three functions. **Split & Invite** ranks second on impact because it is the only proposal here that changes the acquisition equation structurally rather than incrementally.

What would make me abandon each idea

Every MVP in Section 6 carries an explicit kill criterion, and one of them — the Trust Score deposit waiver — can be killed for free, before a line of code is written, by back-testing whether past guest ratings actually predict future damage. If that correlation is weak, the loop is invalid and the work stops there. Designing the cheapest possible route to a negative result is, I would argue, the more important half of MVP design.

SECTION 2

Context & diagnosis

A growth strategy that does not begin with a diagnosis is a list of features. This section establishes what is actually constraining Airbnb in India, using the strongest evidence I could source, and states plainly where I am estimating rather than citing.

2.1 What Airbnb is trying to be

Airbnb’s stated ambition is an end-to-end travel platform spanning accommodation, local experiences and trip logistics, differentiated by personal connection and community between hosts and guests rather than impersonal hotel stays. Its 2025 and 2026 releases show that ambition being executed: Experiences relaunched with a \$200m investment, Services added (airport pickups, luggage storage, car rental, grocery), boutique hotels onboarded in 20 cities, and social features — shared itineraries, a Connections feed showing friends’ trips — shipped in the 2026 Summer Release [R6, R7].

That matters for this document in two ways. First, the scaffolding for several proposals below already exists globally and needs localisation rather than invention. Second, the community positioning is a real constraint: it rules out growth tactics that would work fine for a commodity OTA.

2.2 The evidence base

FINDING	DETAIL	SOURCE
India is overwhelmingly a domestic-travel market for Airbnb	Domestic travellers were 91% of Airbnb’s Indian guests in 2024, up from 79% in 2019, driven by younger guests. India was among Airbnb’s fastest-growing markets globally in 2024.	R1
Airbnb is not price-competitive on commission, and direct booking is free	Indicative Indian host commissions: Airbnb ~15.5% host-only; Booking.com 15–20%; MakeMyTrip/Goibibo 18–25%; Agoda 15–22%; OYO 20–25%+. Direct booking costs 0% and settles instantly over UPI.	R2
Hosts are financially motivated to disintermediate	Typical Indian homestays are estimated to lose ₹3–8 lakh a year to OTA commissions. Airbnb pays out fastest (24h after check-in) — its genuine advantage on this axis.	R2
Guest photos in reviews are a live competitive gap	Booking.com has offered guest photos within reviews since 2016. Vrbo launched them in October 2025, showing them in reviews, the main gallery, and alongside an AI review summary. Airbnb does not offer them.	R3
Referral incentives have already been tried and retired	Airbnb’s Associates (affiliate) programme closed on 31 March 2021 with 30 days’ notice. Any referral proposal must therefore explain why it would not fail the same way.	R4
Guest-generated photos have measurable effects	Peer-reviewed work on guest-generated photos on Airbnb exists in <i>Annals of Tourism Research</i> (2025), supporting the mechanism behind Section 6’s MVP 1.	R5

2.3 The four constraints

Constraint 1 — Trust is the conversion bottleneck

India’s alternative-accommodation supply is largely unprofessional: individually owned, inconsistently maintained, and photographed once at listing time and rarely again. The result is the objection that dominates Indian guest hesitation — “will the place actually look like the pictures?” — compounded by anxiety about last-

minute host cancellation and, for solo and women travellers, safety. Every one of these is a trust problem the product can address, and each one currently converts into either an abandoned search or a support ticket.

Constraint 2 — Frequency, not just acquisition

Indian leisure travel clusters into long weekends, festivals and weddings, producing roughly one to two trips per user per year for most travellers. At that frequency, retention mechanics that work in high-frequency categories do not apply, and customer lifetime value stays low regardless of how efficient acquisition becomes. Raising trips per user per year is therefore a first-order growth lever, not a secondary one — and it is the lever Airbnb India appears to invest in least.

Constraint 3 — Supply is thin and leaks

Outside a handful of leisure markets — Goa, Himachal, Uttarakhand, Kerala, Rajasthan — depth of quality supply is limited, and the supply that exists is actively incentivised to leave. A host losing lakhs a year in commission and able to settle instantly by UPI has both the motive and the means to take repeat guests direct. Airbnb cannot police its way out of this; it has to make staying on-platform the cheaper option.

Constraint 4 — The take rate does not survive the arithmetic

THE CORE ECONOMIC PROBLEM, IN ONE LINE

A ₹4,000-per-night, two-night booking at a 15.5% take rate yields roughly **₹1,240 of revenue** — against real payments, support and trust costs, in a market where MakeMyTrip and Goibibo charge 18–25% with far deeper inventory and where going direct costs the host nothing at all. Airbnb is simultaneously **too expensive for the host** and **too small in absolute revenue for the platform**. Raising the take rate accelerates the leakage; holding it flat concedes the market.

2.4 The positioning constraint

Airbnb's differentiation is human connection. That is not decoration — it is a hard filter on what may be proposed. Coupon-led acquisition, cash-for-generic-referral schemes and a price-parity war against MakeMyTrip are all available, all effective in the short term, and all off-strategy, because they compete on exactly the axis where Airbnb is weakest and its competitors are strongest. Section 5.4 applies an explicit test to every loop proposed here: *would this work identically for Booking.com?* If yes, it is not a defensible loop, and I say so.

2.5 The North Star metric

PROPOSED NORTH STAR: VERIFIED TRIPS COMPLETED PER ACTIVE TRAVELLER PER YEAR

Why this one. It contains all four constraints in a single number. *Verified* forces trust and completion quality rather than raw bookings. *Trips completed* excludes cancellations, which is where Indian trust failures actually show up. *Per active traveller per year* makes frequency — the metric most easily ignored — impossible to ignore. A team optimising this number cannot hit it through discounting alone.

SECTION 3

Product audit — what exists vs. what is missing

Proposing a feature that already ships is the most common failure mode in growth strategy work. This section records what I checked before proposing, and how each idea was repositioned once the audit came back.

CANDIDATE IDEA	STATUS ON AIRBNB TODAY	VERDICT	REPOSITIONED AS
Guest photos in reviews	Not available. Competitors have shipped it — Booking.com since 2016, Vrbo since October 2025 [R3].	GAP — BUILD	Loop D — Proof of Stay. Kept as originally conceived; it is the strongest single proposal in this document.
Two-way host ↔ guest reviews	Already exists , and has since launch. Double-blind, 14-day window; hosts rate guests on cleanliness, communication and house rules [R8].	EXISTS	Loop E — Trust Score. The feature is not missing; the <i>incentive</i> is. The guest rating is invisible and carries no consequence, so it changes no behaviour.
Social referral rewards	Tried and retired. The Associates programme closed 31 March 2021 [R4]. Generic invite links attract low-intent users and are easily gamed.	RETRY WITH CHANGES	Loop F — Host-Funded Advocacy. Two changes: listing-specific rather than generic, and host-funded rather than platform-funded.
Guest-curated recommendations	Partially exists — Wishlists, Guest Favorites (2023), and the Connections feed in the 2026 Summer Release [R7, R9].	PARTIAL	Loop G — Guest Curator. Curation exists; creator <i>economics</i> do not. Publishing and attribution turn a bookmark into a distribution channel.
AI customer support	Exists — resolving ~94% of English-language enquiries, with AI support in 11 languages announced for 2026 [R6, R7].	PARTIAL	Initiative 1.5. The gap is channel and language depth for India: WhatsApp-native, Hindi and regional languages.
Trip services / logistics	Exists globally — airport pickups, luggage storage, car rental, grocery delivery via partners [R7].	PARTIAL	Revenue Layer 2. The India localisation is different: cook-at-home and drivers matter far more than grocery delivery.

WHAT THE AUDIT CHANGES ABOUT THE ARGUMENT

Two of six candidate ideas needed repositioning from “new feature” to “existing feature, broken incentive.” That reframe is more valuable than the original ideas were, because a feature that already exists has already absorbed its build cost — making the incentive work is cheaper than building something new, and the counterfactual is much easier to measure. The Trust Score proposal in particular is now a *configuration and incentive* change on top of a decade-old data asset, rather than a new system.

SECTION 4

PLG initiatives for cross-functional growth

The brief asks how product-led growth helps *cross-functions* — Sales, Marketing, Support, Operations — succeed at driving growth. That is a subtly different question from “how does Airbnb do PLG,” and the difference matters.

4.1 Framing: product-led sales, support and operations

Classical PLG asks how the product acquires and expands users without a salesperson. The question here is the inverse and more interesting one: **how does the product make the humans in the organisation more effective?** In a low-ABV market like India this is not a nice-to-have. At roughly ₹1,240 of revenue per booking, any function whose unit cost scales linearly with bookings is structurally unaffordable. The product must either absorb that work or make the remaining human effort dramatically better targeted.

ON “SALES”

Airbnb has no classical sales function in India in the enterprise sense — there is no quota-carrying rep closing deals. The equivalent motion is **supply acquisition**: persuading a property owner to list and to complete listing. That is the motion initiatives 1.1 and 1.2 address, and it is why they are grouped under “Supply Operations & Sales” below. Where a genuine sales motion does exist — professional property managers and multi-property hosts — the Co-Host Marketplace is the product surface that serves it.

Three mechanisms recur across every initiative below:

- **Absorption** — the product does the task outright (AI listing creation, self-serve resolution).
- **Targeting** — the product decides where scarce human effort is spent, using in-product signals rather than a static list.
- **Manufacture** — the product generates an asset another function would otherwise buy (marketing creative, SEO inventory, trust signals).

4.2 The initiatives

Supply Operations & Sales

INITIATIVE	PROBLEM TODAY & PRODUCT-LED MECHANISM	FUNCTION KPI
1.1 Host Activation Copilot TARGETING	Prospective hosts abandon listing creation at predictable points — photography, pricing, and description. Ops chases them blind, in list order, with no signal about who is actually winnable. The product should do two things: an AI listing builder (photo quality scoring, auto-generated copy, comparable-based price suggestion) that removes the drop-off causes outright, and a Host Readiness Score that feeds a signal-ranked ops queue, so a human is dispatched only where the model predicts human contact changes the outcome.	Listings published per ops FTE; time-to-first-booking
1.2 Co-Host Marketplace ABSORPTION	Supply acquisition in India is a one-to-one manual motion that does not scale to tier-2 and tier-3 cities. But India already has an informal layer of caretakers, managers and family members who operate properties on owners' behalf. Productising that layer — verified co-hosts who can list, service and manage on behalf of owners — converts a headcount-bound activity into a marketplace that recruits supply on Airbnb's behalf.	Listings sourced via co-hosts; supply CAC

Marketing

INITIATIVE	PROBLEM TODAY & PRODUCT-LED MECHANISM	FUNCTION KPI
1.3 Programmatic Trust Pages MANUFACTURE	Acquisition leans on paid channels, while the long tail of high-intent Indian queries — “ <i>villas for a 12-person wedding near Lonavala</i> ”, “ <i>pet-friendly homestay Coorg</i> ” — is under-served. Listings, structured amenities, reviews and guest photos should compose automatically into intent-matched landing pages. The inventory is already in the database; only the composition layer is missing.	Organic sessions; blended CAC
1.4 Guest-Generated Creative Engine MANUFACTURE	Creative is commissioned, expensive and slow, and studio photography is exactly what Indian guests distrust. Structured post-stay capture — the same flow as MVP 1 — yields a continuous stream of authentic, geo-tagged, permissioned imagery. This initiative is a direct output of Loop D , which is what makes Proof of Stay disproportionately valuable: one feature, three functions.	Creative cost per acquisition; usable assets/month

Customer Support

INITIATIVE	PROBLEM TODAY & PRODUCT-LED MECHANISM	FUNCTION KPI
1.5 WhatsApp-native multilingual resolution ABSORPTION	Airbnb’s AI support resolves ~94% of English enquiries [R6], but India is neither English-first nor app-first. Extending resolution to Hindi and the major regional languages, delivered over WhatsApp , does two things at once: it collapses cost-to-serve on a low-ABV booking, and it places Airbnb inside the channel where Indian trip planning actually happens — making a cost centre into a re-engagement surface.	Contact rate per 100 bookings; cost per booking
1.6 Ticket-to-Roadmap loop TARGETING	Every ticket is a product failure that has already happened. A fixed contact taxonomy, reviewed monthly, with the top driver becoming a committed product fix owned by a named PM, converts support from a queue into the highest-signal research function in the company.	% of contacts from the top-3 drivers (declining)

Operations, Trust & Safety

INITIATIVE	PROBLEM TODAY & PRODUCT-LED MECHANISM	FUNCTION KPI
1.7 Verified India Layer ABSORPTION	Trust & Safety casework is manual, and the underlying guest anxieties — listing accuracy, cancellation risk, personal safety — suppress conversion long before they generate a case. An opt-in verification layer using India’s existing digital identity rails, OTP or smart-lock check-in, assisted foreigner registration (Form C) for inbound hosts, and an explicit cancellation guarantee converts casework into a <i>conversion feature</i> . This is the data source that makes Loop E’s Trust Score possible.	Manual T&S cases per 1,000 bookings; search → book conversion
1.8 Shared growth data layer TARGETING	The unglamorous prerequisite. One event schema and one set of agreed metric definitions, so Supply Ops, Marketing, Support and Product act on the same signals rather than four reconciliations of the same question. Without it, initiatives 1.1 and 1.6 cannot function, because neither has a trustworthy signal to rank on.	Time from signal to action

4.3 Why these are PLG and not just operational improvements

THE TEST

An operational improvement makes a team faster at the work it already does. A **product-led** initiative removes the work, or makes the product decide where the work goes. Initiatives 1.1, 1.2, 1.5 and 1.7 remove work. Initiatives 1.3, 1.4 and 1.6 make the product manufacture something a function would otherwise buy or discover manually. Initiative 1.8 is the substrate that lets the other seven use in-product behaviour as their trigger. None of them add headcount to grow output — which is the actual definition being tested.

4.4 How Section 4 connects to Section 5

These initiatives are not independent of the growth loops that follow. Two dependencies are worth naming explicitly, because they are what make the strategy a system rather than a list:

- **Initiative 1.4 is powered by Loop D.** The Guest-Generated Creative Engine has no input until Proof of Stay ships. Build the loop, and the marketing initiative becomes free.
- **Loop E is powered by Initiative 1.7.** The Trust Score needs a verified identity substrate to be meaningful; without it, the score is gameable and the deposit waiver becomes a fraud vector.

This is the practical answer to the brief's first requirement. Cross-functional growth does not come from a product team handing finished features to other functions. It comes from designing loops whose *by-products* are the raw material those functions need — trust signals for Ops, creative for Marketing, deflection for Support, qualified leads for Sales.

SECTION 5

Growth loops & the India flywheel

A funnel is linear and leaks: it must be refilled from the top with spend, forever. A loop's output becomes its own next input, so each turn makes the next one cheaper. This section defines the flywheel Airbnb India should be turning, then the eight loops that feed it.

5.1 The India flywheel




Figure 1. The flywheel Airbnb India should be turning. India is weakest on trust and frequency, which is why the loops below are weighted toward those two inputs rather than toward selection.

The flywheel itself is not novel — it is roughly the marketplace flywheel Airbnb has always run. What matters is **where energy is injected**. In a mature Western market, selection is the binding input and supply acquisition is the lever. In India, selection is not the first thing that breaks: a guest who does not trust the photographs never gets far enough to notice that selection is thin. Trust gates the entire cycle, and frequency determines how many times per year a user re-enters it.

5.2 Three engines, eight loops

Grouping the loops into three engines — each feeding one input of the flywheel — makes the portfolio legible and exposes where it is thin.

🛡️ **Engine 1 — Trust** → feeds booking confidence

Loop D — Proof of Stay · Loop E — Two-Way Trust Score

👥 **Engine 2 — Social distribution** → feeds acquisition

Loop B — Split & Invite · Loop F — Host-Funded Advocacy · Loop G — Guest Curator




Engine 3 — Supply & frequency → feeds selection and repeat

Loop A — Host Someone Like You · Loop C — Local Host Loop · Loop H — Weekend Radius

5.3 The four priority loops in detail

Loop D — “Proof of Stay” **HIGHEST PRIORITY**




Guests attach verified photographs to reviews. Those photographs appear in the review, and — in a later iteration — in the listing gallery alongside the host’s own. Because they come from completed, verified stays, they cannot be staged by the host.

Why it is the priority. It is a confirmed competitive gap: Booking.com has had guest photos since 2016 and Vrbo shipped them in October 2025, while Airbnb has not [R3]. It attacks the top Indian guest objection directly. And it compounds across three functions at once — trust asset, marketing creative, SEO content.

THE OBSTACLE, NAMED AND SOLVED

Airbnb has plausibly avoided this feature to protect hosts from unflattering images and to keep listing pages visually controlled. That is a real concern and a proposal that ignores it is not credible. The mitigations: moderation before publication, a host right-of-reply on every photo review, photos restricted to verified completed stays, and — in v1 — display confined to the reviews section rather than the main gallery. Host delisting rate is therefore a hard guardrail in MVP 1, not an afterthought.


Loop E — “Two-Way Trust Score”



Reciprocal reviews are not a new idea for Airbnb — they are one of its oldest. Hosts already rate guests on cleanliness, communication and adherence to house rules, double-blind, within a 14-day window [R8]. **The defect is that the guest’s resulting rating is nearly invisible and carries no consequence**, so it does not alter anyone’s behaviour and the loop never closes.

Attaching real benefits closes it: a waived or reduced security deposit, Instant Book access, priority support, and eventually a reduced service fee. In India this is unusually well-targeted, because deposit friction and host anxiety about property damage are both live suppressors — of guest conversion on one side and host willingness to accept bookings on the other. One mechanism relieves both.


Loop F — “Host-Funded Advocacy”



Airbnb has run a referral programme before and retired it [R4]. Any proposal to run one again has to answer why it would not fail identically. The honest reading of that failure: generic invite links attract low-intent signups, are trivially gamed, carry no quality signal, and are funded entirely from the platform's balance sheet. Two design changes address each of those causes.

1. **Listing-specific, not generic.** A guest shares the actual property they stayed in, with an attributed link. Intent is higher, gaming is harder, and the share doubles as authentic marketing for that host.
2. **Host-funded, not platform-funded.** This is the economic flip that matters. The reward comes from margin the host is *already losing* to commission. Airbnb's marginal cost approaches zero, so the loop sustains at scale — and the host gains an acquisition channel materially cheaper than the 15.5% commission, which attacks direct-booking leakage by making the platform the cheaper option rather than by policing hosts.

Loop B — “Split & Invite”



India travels in groups — friends, extended families, wedding parties. Today that produces exactly **one** Airbnb account: the organiser pays, and the other five to seven settle up privately over UPI or Splitwise, never touching the product. Making the booking a **multiplayer object** converts a cultural default into a viral coefficient at near-zero customer acquisition cost.

The scaffolding partly exists — the 2026 Summer Release ships shared itineraries and a Connections feed [R7]. The India-specific unlock is **UPI split payment**, which is what makes per-person settlement natural rather than awkward. This is the proposal here that most changes the structure of acquisition rather than its efficiency.

5.4 Positioning check — would this work for Booking.com?

Airbnb's stated edge is human connection, not commodity rooms. A loop that a pure OTA could copy wholesale is not a defensible loop. Applying that test honestly, including where it is unflattering:

LOOP	DIFFERENTIATED?	REASONING
A — Host Someone Like You	YES	Depends entirely on the host – guest relationship and a mentor drawn from the host community. A hotel chain has no equivalent.
B — Split & Invite	YES	Group travel plus the community/shared-trip framing; the social layer is Airbnb-native.

C — Local Host Loop	YES	A community of local experience hosts <i>is</i> the asset. Not replicable by inventory aggregation.
D — Proof of Stay	MECHANIC IS GENERIC	Booking.com has had it since 2016 [R3]. The argument here is not differentiation but gap-closing — a different and equally valid case, stated as such.
E — Two-Way Trust Score	YES	Reciprocal reputation only makes sense where a person hosts a person. A hotel does not review its guests.
F — Host-Funded Advocacy	YES	Host co-funding is only possible with independent hosts who control their own margin. A chain cannot devolve this.
G — Guest Curator	YES	Community-to-community recommendation, consistent with the positioning.
H — Weekend Radius	PARTIALLY	The mechanic is generic; differentiation comes only from drive-to supply an OTA does not hold.

EXPLICITLY REJECTED AS OFF-POSITIONING

Cash-for-generic-referral schemes; coupon-led acquisition; a displayed-price parity war with MakeMyTrip. All three would move short-term numbers. All three compete on precisely the axis where Airbnb is weakest and its competitors strongest, and all three erode the community positioning that makes loops A, B, C, E, F and G possible in the first place. They are excluded deliberately, not overlooked.

5.5 Loops requiring MVPs

Four loops involve substantial new features and therefore require MVP designs: **D (Proof of Stay)**, **F (Host-Funded Advocacy)**, **B (Split & Invite)** and **E (Trust Score benefits)**. These are designed in Section 6. Loops A, C, G and H are lower-priority and are treated in the roadmap in Section 8 rather than given full MVP designs.

SECTION 6

MVP designs

Each MVP below follows the same structure: a falsifiable hypothesis, an assumption map ranked by risk, the cheapest prototype that attacks the riskiest assumption, layered success metrics, and an explicit kill criterion. All four are implemented in the companion prototype; the screenshots are taken from it.

THE DISCIPLINE BEING APPLIED

An MVP is not a small version of the product. It is **the cheapest experiment that can produce a decisive negative result**. Accordingly, the riskiest assumption in each design below is almost never technical — it is behavioural or economic. MVP 4 is the clearest illustration: it can be killed by a retrospective data analysis costing nothing, before any code is written.

MVP 1 · LOOP D · ENGINE: TRUST

Proof of Stay — verified guest photos in reviews

*We believe that letting guests attach verified photos to their reviews will increase search → book conversion on listings with guest photos by **at least 5% relative**, and that at least **15% of reviews** will include a photo, within 8 weeks, across 3 Indian cities.*

Assumptions, ranked by risk

- **Guests will upload photos without a large incentive.** Riskiest. This is a content-supply problem, and if it fails, nothing downstream matters. Attacked first by placing the ask at the moment of peak goodwill.
- **Hosts will not revolt or delist over unflattering photos.** The most likely reason Airbnb has not already shipped this. Mitigated by moderation, right-of-reply, and reviews-only placement in v1.
- **Photos increase trust rather than choice-paralysis.** Plausible but unproven; measured by conversion, not by sentiment.
- **Moderation is tractable at Indian volumes.** Low risk; manual in the MVP, automatable later.

Prototype — deliberately narrow

Three cities only. The upload prompt fires immediately after a 5★ review is submitted — the point of maximum goodwill. Photos appear in the **reviews section only**, not the main gallery, which caps host backlash in v1. Moderation is performed manually by an ops associate; no machine-learning classifier is built. Hosts receive a right-of-reply on every photo review. Estimated effort: ~3 person-months.




Figure 2. Control vs. variant on the same listing, captured from the prototype's demo-control panel. The variant adds a verified guest-photo strip, a “with photos” review filter, and photos attached to individual reviews.




Figure 3. The post-stay capture flow — step 3 of 4. Copy explicitly frames the upload as helping the next guest, and the “Verified stay” label is the trust signal being manufactured.

Success metrics

LAYER	METRIC	TARGET
Leading (days)	% of submitted reviews including ≥1 photo	≥ 15%
Primary (weeks)	Search → book conversion on listings with guest photos vs. matched control	≥ +5% relative
Guardrail	Host delisting / churn rate in the test cohort	< +1pp

Kill criteria: photo upload rate below 5% of reviews, or host delisting rate up more than 1pp in the test cohort. Either result invalidates the loop — the first because the content supply does not exist, the second because the cost is borne by the wrong side of the marketplace.

MVP 2 · LOOP F · ENGINE: SOCIAL DISTRIBUTION

Host-Funded Advocacy

*We believe that letting hosts fund a shareable, listing-specific reward will produce incremental bookings at a blended acquisition cost **below the 15.5% commission benchmark**, with at least **20% of eligible hosts opting in**, within one quarter.*

Assumptions, ranked by risk

- **Hosts will voluntarily give up margin** for a booking they might have received anyway. Riskiest, and it is economic rather than technical. If hosts will not fund it, the model reverts to platform-funded — exactly the design that was retired in 2021.
- **Guests will share a specific property publicly.** Sharing a listing reveals where you stayed and holiday plans; some guests will not.
- **Shares produce incremental bookings, not just attribution of bookings that would have happened.** Requires a holdout to measure honestly.
- **Attribution and reward settlement are technically feasible.** Standard tracked-link mechanics.

Prototype — zero engineering

Recruit **30 hosts across 2 cities** by hand. Each receives a trackable link and a co-funded offer configured manually in a spreadsheet; ops reconciles bookings and issues credits by hand. This is a full Wizard-of-Oz build: nothing is automated, and the entire economic question — will hosts pay, and does it beat commission — is answerable without shipping software. Estimated effort: ~3 person-months, most of it operational rather than engineering.




Figure 4. The host-side configuration screen from the prototype. The reward slider recalculates the comparison against commission live. This panel is the entire argument for the loop: at a ₹500 reward on a ₹28,500 booking, the host's effective acquisition cost is 1.8% against a 15.5% commission benchmark.

Success metrics

LAYER	METRIC	TARGET
Leading	Host opt-in rate; shares per completed stay	≥ 20%; ≥ 0.3
Primary	Blended cost per acquired booking vs. the 15.5% commission benchmark	< 15.5%
Guardrail	Self-referred / fraudulent booking rate	< 1%
Guardrail	Guest NPS; unsubscribe rate on share prompts	Flat

Kill criteria: host opt-in below 10%, or cost per acquired booking at or above the commission benchmark. The second is decisive — if the channel is not cheaper than commission, the host has no reason to fund it and the loop has no engine.

MVP 3 · LOOP B · ENGINE: SOCIAL DISTRIBUTION

Split & Invite — UPI group booking

*We believe that letting a booking organiser split payment across co-travellers via UPI will produce **at least 1.5 newly activated accounts per group booking**, within 8 weeks, across 3 cities.*

Assumptions, ranked by risk

- **Organisers want to formalise splitting on-platform** rather than settling privately. Riskiest, and purely behavioural. Splitting already works fine on UPI and Splitwise; the question is whether doing it inside the booking is worth the friction of adding people.

- **Co-travellers will complete an account** to pay their share, rather than simply paying the organiser back.
- **Splitting does not depress checkout completion.** Adding steps to checkout is inherently risky; hence the hard guardrail.
- **UPI collect/autopay integration is feasible.** Mature rails, well-documented.

Prototype — staged, near-zero cost at stage 1

Stage 1 (2 weeks) — painted door. A “Split this booking” button appears at checkout in 3 cities. Clicking it returns “coming soon — notify me.” This measures intent alone and costs essentially nothing. **Stage 2 (6 weeks) — concierge.** For 100 real group bookings, an ops associate manually creates the shared itinerary and issues UPI collect links. No payment infrastructure is built until both stages pass.




Figure 5. The confirmation state from the prototype. The banner quantifying newly activated accounts is the loop's payoff made visible — it is the moment a single booking becomes five users.

Success metrics

LAYER	METRIC	TARGET
Leading	Painted-door click-through rate at checkout	≥ 8%
Primary	Newly activated accounts per group booking	≥ 1.5

Primary	90-day booking rate of invited co-travellers	≥ 6%
Guardrail	Checkout completion rate	> -2pp

Kill criteria: painted-door CTR below 3%, or fewer than 0.5 activated accounts per booking in the concierge stage. Note that stage 1 alone can kill the idea for roughly two weeks of work — which is the point of staging it.

MVP 4 · LOOP E · ENGINE: TRUST

Trust Score → security deposit waiver

*We believe that waiving the security deposit for guests above a Trust Score threshold will increase booking conversion by **at least 3%** without raising the damage-claim rate by more than **0.5pp**.*

Assumptions, ranked by risk

- **Past guest ratings actually predict future damage.** Riskiest by far — and testable for free. If the correlation is weak, the entire loop is invalid and no code should be written.
- **Deposit friction meaningfully suppresses conversion** at Indian price points.
- **Hosts will accept the platform underwriting the waived deposit.**
- **The score can be computed from existing review data.** The data has existed for over a decade.

Prototype — the cheapest kill in the set

STEP 0 COSTS NOTHING AND CAN END THE PROJECT

Before any build: **back-test on historical bookings.** Take past guest ratings and test whether they predict subsequent damage claims. This is a retrospective data analysis requiring no engineering, no users and no risk. If there is no predictive signal, the feature dies for the price of an analyst's week. Only if the signal holds does a manual whitelist pilot with 200 high-score guests follow.




Figure 6. The moment the Trust Score becomes consequential. A ₹10,000 deposit struck through to ₹0 cuts the total by nearly a quarter (₹42,490 → ₹32,490) — this is what turns an invisible rating into a behavioural incentive.

Success metrics

LAYER	METRIC	TARGET
Step 0	Correlation between historical guest rating and subsequent damage claims	Statistically significant
Leading	% of active guests above the qualifying threshold	Sizing only
Primary	Booking conversion lift in the waived cohort vs. holdout	≥ +3%
Guardrail	Damage claim rate; claim value per booking	< +0.5pp
Secondary	Host Instant Book adoption	Increasing

Kill criteria: no predictive signal in the historical back-test. This is the only kill criterion in the document that can be evaluated before the product is built, which makes MVP 4 the first piece of work I would schedule despite it not being the highest-priority feature.

6.5 Experiment dashboard

The prototype includes a mock analytics view rendering each hypothesis against its live metrics and kill criteria — the instrumentation as it would appear at week 8 of the four experiments running in parallel.




Figure 7. The experiment dashboard from the prototype, showing MVP 1 at week 8. Each hypothesis appears above its own leading, primary and guardrail metrics, with the kill criterion stated and its status shown. Metrics are illustrative.

6.6 What links the four MVPs

The four designs are not independent bets. MVP 4's back-test is scheduled first because it is free and can produce a negative result before anything is built. MVP 1 is scheduled next because its output — verified photographs — is the quality signal that Section 7's Earn-Out Commission later needs in order to tier hosts credibly. MVP 3 is staged so that two weeks of work can kill it. And MVP 2 is deliberately built with no engineering at all, because its riskiest assumption is economic rather than technical.

Read together, the four form a sequence in which each experiment is cheaper than the feature it authorises, and every one of them can end in a documented “no.” That is the property I would most want a growth portfolio to have: not that every bet succeeds, but that the failures are cheap and legible.

SECTION 7

Revenue model for India

The strategic claim of this section is a single sentence: **take less on the stay, earn more on the trip**. What follows is the case for that, the model, the arithmetic, and the sensitivity.

7.1 Why the current model does not work in India

Airbnb monetises through a percentage take rate on booking value — approximately 15.5% under the host-only fee structure that became default for most hosts in late 2025 [R2]. That model is elegant in high-ABV markets. In India it fails on three counts at once.

- **The absolute number is too small.** A ₹4,000-per-night, two-night booking yields roughly ₹1,240 of gross revenue, against real payments, support and trust costs.
- **It is not competitive for the host.** MakeMyTrip and Goibibo charge 18–25% but offer far deeper inventory and demand [R2]. Airbnb charges less than the incumbents but delivers less volume — the worst position on that trade-off curve.
- **The alternative is free.** Direct booking costs the host 0% and settles instantly over UPI [R2]. With estimated annual commission losses of ₹3–8 lakh per homestay, hosts have both a strong motive and a frictionless mechanism to disintermediate.

A NOTE ON WHICH FEE STRUCTURE IS MODELLED

Airbnb operates two fee structures: a **host-only** fee of roughly 15.5%, and a **split fee** where the host pays ~3% and the guest pays a further 14–16% [R2]. The unit economics in §7.3 assume the **host-only** structure throughout, because it books less revenue per booking and is therefore the conservative basis for arguing that contribution can rise while the take rate falls. Two places in this document depict the split-fee structure instead, and it is worth flagging rather than glossing: the prototype checkout in Figure 6 shows a guest service fee, and Layer 3a’s “zero guest service fee” benefit is only meaningful to guests on the split structure. Under host-only, that subscription benefit would instead be a service-fee rebate or a credit — the mechanism changes, the ARR argument does not.

WHY “JUST RAISE THE TAKE RATE” IS THE WRONG ANSWER

Raising the take rate increases revenue per retained booking while accelerating the rate at which hosts leave — and in a market where the substitute costs zero and settles instantly, elasticity is high. The model below therefore moves in the opposite direction on the headline rate and recovers the difference elsewhere.

7.2 The proposed model — five layers

L	LAYER	MECHANISM & INDIA RATIONALE
1	Earn-Out Commission RETENTION	A take rate that <i>declines</i> with earned quality and loyalty: 15.5% → 12% → 9% as a host clears verified-quality and repeat-guest thresholds. This inverts the usual dynamic — the platform becomes cheaper the more loyal the host is, attacking leakage by making staying the rational choice rather than by enforcement. Quality is measured using assets this document already creates: photo-verified reviews (Loop D) and Trust Score data (Loop E).

2	Trip Services attach GROWTH ENGINE	Monetise the trip, not the room: intercity cabs and airport transfers, cook/chef-at-home , housekeeping, experiences, cancellation protection, luggage. This is incremental spend that never touches the price-sensitive room rate. Airbnb's 2026 Services rollout [R7] is the global template; the India localisation is materially different — cooks and drivers, not grocery delivery.
3a	Guest subscription “Weekend Club” ₹999–1,499/yr	Zero guest service fee on domestic stays, free cancellation, priority Hindi support, and automatic Trust Score qualification. Converts a per-transaction fee into predictable annual revenue <i>and</i> raises frequency through sunk-cost commitment — attacking the fee-sensitivity and frequency problems with one instrument.
3b	Host subscription “Host Pro” ₹499–999/mo	Dynamic pricing, co-host marketplace access, channel manager, GST invoicing , WhatsApp guest comms, and advocacy campaign tools (Loop F). Critically, this monetises hosts <i>on nights they sell elsewhere</i> — revenue from supply that Booking.com and MakeMyTrip currently capture for free.
4	Demand-side advertising CAPPED / EXPERIMENTAL	CPC sponsored placement for hosts in low season. High margin, but it creates a direct conflict between ranking integrity and revenue — precisely the dynamic that erodes trust on incumbent OTAs. Recommended only as a capped experiment with a hard guardrail on booking satisfaction, and named here as a risk rather than presented as free money.
5	Payments	UPI and UPI Autopay, pay-in-parts, no-cost EMI on high-value villa and wedding bookings, and instant refunds. Payment friction is a conversion tax in India; instant refunds additionally attack the cancellation-distrust problem identified in Section 2.

7.3 Unit economics — illustrative

The model below is built on a representative booking of **₹4,000 per night × 2 nights = ₹8,000 GMV**. All figures are illustrative and constructed from stated assumptions, not from Airbnb's internal data; the purpose is to test whether the *structure* can work, not to forecast a result.

PER BOOKING (₹8,000 GMV)	TODAY	PROPOSED	ASSUMPTION
Commission revenue	1,240	1,040	15.5% → blended 13% under Earn-Out
Trip services attach	—	90	30% attach × ₹2,000 spend × 15% take
Guest subscription (amortised)	—	60	12% penetration × ₹1,199/yr ÷ 2.4 trips
Host subscription (amortised)	—	14	8% of hosts × ₹699/mo ÷ 4 bookings/mo
Total revenue	1,240	1,204	
Payments cost	(160)	(96)	2.0% → 1.2% on UPI rails
Support cost	(120)	(90)	WhatsApp AI deflection (Initiative 1.5)
Trust & safety cost	(80)	(50)	Trust Score reduces claims & casework
Contribution per booking	880	968	+10%

THE RESULT, AND WHY IT IS CONSERVATIVE

The headline take rate **falls from 15.5% to a blended 13%**, and contribution per booking still **risks about 10%**. And this model excludes the two largest effects of the strategy: the additional volume from reduced direct-booking leakage (Layer 1 and Loop F), and the additional trips per user from Loop B and Loop H. Including those would improve the result materially — which is why the table above should be read as the conservative case rather than the expected one.

Sensitivity — contribution per booking (₹)

The model's single most uncertain input is trip-services attach rate. Testing it against subscription penetration:

ATTACH RATE ↓ / GUEST SUBSCRIPTION PENETRATION →	6%	12%	20%
15% attach (pessimistic)	893	923	963
30% attach (base)	938	968	1,008
45% attach (optimistic)	983	1,013	1,053

Every cell in the grid — including the most pessimistic corner at ₹893 — exceeds today's ₹880. The model is therefore robust to being substantially wrong about its riskiest input — which is the property one actually wants from a revenue proposal.

7.4 Regulatory and practical considerations

- **GST and e-commerce operator obligations.** As an ECO, the platform carries TCS collection and reporting duties on supplier transactions, and short-term accommodation attracts GST at rates that vary by tariff. Any subscription product must be priced GST-inclusive to avoid a displayed-price disadvantage.
- **All-inclusive displayed pricing.** Indian users compare on the first visible number. Fees revealed late are both a conversion killer and a trust cost — which matters doubly given Section 2's diagnosis.
- **Host tax literacy.** GST invoicing inside Host Pro is not a minor feature; for many small hosts it is the single most valuable thing the subscription could offer, and it deepens platform dependence in a way commission never can.

7.5 Validating the revenue model cheaply

MINI-MVP: FAKE-DOOR PAYWALL FOR WEEKEND CLUB

Hypothesis: at least 8% of active Indian guests will express purchase intent for a ₹1,199/year membership. **Prototype:** a pricing page shown in one city, ending at “join the waitlist” — no billing infrastructure is built. **Primary metric:** click-to-intent conversion. **Kill:** intent below 3%. Willingness to pay is measured *before* any subscription system exists, applying exactly the same discipline as Section 6.

SECTION 8

Prioritisation & roadmap

8.1 RICE scoring

Reach is indexed per quarter (100 = all active Indian users). Impact uses the standard 0.25–3 scale. Confidence reflects evidence strength — where a claim rests on a verified competitive gap or existing data, confidence is high; where it rests on untested behavioural assumptions, it is low. Effort is in person-months.

INITIATIVE	REACH	IMPACT	CONF.	EFFORT	RICE	RANK
Proof of Stay (MVP 1)	100	2.0	0.8	3	53.3	1
Weekend Radius (Loop H)	80	1.5	0.5	2	30.0	2
Trust Score benefits (MVP 4)	70	2.0	0.6	3	28.0	3
WhatsApp support (1.5)	100	1.0	0.9	4	22.5	4
Split & Invite (MVP 3)	40	3.0	0.7	4	21.0	5
Host-Funded Advocacy (MVP 2)	30	2.0	0.5	3	10.0	6
Trip Services attach (Rev L2)	60	2.5	0.5	8	9.4	7
Weekend Club (Rev L3a)	50	2.0	0.4	5	8.0	8
Host Activation Copilot (1.1)	20	2.0	0.7	5	5.6	9
Guest Curator (Loop G)	25	1.0	0.4	3	3.3	10

WHERE RICE MISLEADS, AND WHY I AM OVERRIDING IT

RICE systematically undervalues **Split & Invite**, because its reach is scored on group bookings only. But it is the sole proposal here that changes the *structure* of acquisition rather than its efficiency — every other initiative makes an existing motion better, while this one makes one booking produce five users. I would therefore fund it above its rank. Conversely, Trip Services attach scores poorly on effort but carries most of the revenue upside in Section 7, so it must start early even though it delivers late. **A score is an input to the decision, not the decision.**

8.2 Roadmap

The sequencing below follows the dependency chain rather than the RICE order. Layer 1 of the revenue model needs a credible quality signal, and that signal does not exist until Proof of Stay and Trust Score have been running long enough to produce one — so attempting the revenue redesign first would mean tiering commission on data that is not yet trustworthy.

HORIZON	SHIP	PROVES
0–3 months Prove trust	MVP 4 Step 0 back-test (week 1 — free, may kill the feature); MVP 1 Proof of Stay in 3 cities; MVP 3 stage 1 painted door; Initiative 1.8 shared data layer.	That trust content can be sourced from guests and that it moves conversion
3–9 months Prove loops	Proof of Stay to national + gallery placement; Trust Score benefits live; MVP 3 concierge → build; MVP 2 host advocacy pilot; WhatsApp multilingual support; Weekend Radius.	That the loops compound and that frequency responds
9–18 months Prove economics	Earn-Out Commission rollout; Trip Services attach (cook, driver, transfers); Weekend Club and Host Pro subscriptions; Co-Host Marketplace; Guest Curator.	That contribution per booking rises while the take rate falls

8.3 What happens if an early experiment fails

A roadmap that only describes the success path is not a plan. Two of the horizon-one experiments can plausibly fail, and each failure has a defined consequence rather than an improvisation.

If the Trust Score back-test finds no predictive signal, Loop E is invalid and the deposit waiver is abandoned in week one. Section 7's Earn-Out Commission then loses one of its two quality inputs and must tier on photo-verified review data alone, which raises the confidence bar on MVP 1 rather than stopping the revenue work.

If Proof of Stay triggers host delisting above the guardrail, the feature does not simply stop — the finding is that the cost lands on the wrong side of the marketplace. The fallback is host-controlled inclusion: guest photos published only on listings whose hosts opt in, which weakens the trust signal but preserves the marketing and SEO by-products in Initiatives 1.3 and 1.4.

In both cases the roadmap continues. That is the practical value of writing kill criteria before the work starts — the team already knows what the next move is when a bet does not pay off.

SECTION 9

Risks, assumptions & limitations

9.1 Principal risks

RISK	DESCRIPTION	MITIGATION
Host backlash on guest photos	The most likely reason this feature does not already exist. Unflattering images could trigger delisting, particularly among the professional supply Airbnb most wants to retain.	Moderation, right-of-reply, verified-stay gating, reviews-only placement in v1, and host delisting rate as a hard kill criterion.
Advocacy repeats the 2021 failure	If hosts will not fund rewards, the programme drifts back to platform funding and fails the same way the Associates programme did.	Host opt-in is the leading metric and the kill criterion. A Wizard-of-Oz pilot answers the economic question before any build.
Trust Score gaming / fraud	A waived deposit is a fraud incentive. Score inflation or account cycling would create real losses.	Dependency on Initiative 1.7's verified identity layer; damage-claim rate as guardrail; staged threshold rollout.
Cannibalisation by Earn-Out Commission	Hosts who would have stayed anyway receive a discount, reducing revenue with no retention benefit.	Tier on <i>earned</i> thresholds only; measure against a holdout of eligible-but-untiered hosts before full rollout.
Ads erode ranking integrity	Layer 4 creates a structural conflict between revenue and result quality — the dynamic that has damaged trust on incumbent OTAs.	Capped, low-season only, with booking satisfaction as a guardrail. Recommended as experimental, and droppable.
Regulatory change	Short-term rental regulation, GST treatment and municipal rules vary by state and are actively evolving.	State-by-state rollout sequencing; compliance tooling in Host Pro converts a risk into a monetisable feature.

9.2 Assumptions this document rests on

- **Commission benchmarks are indicative, not audited.** The 15.5% Airbnb rate and 18–25% MakeMyTrip range come from an industry guide [R2], not from company filings. Directionally reliable; not precise.
- **All unit economics in Section 7 are constructed illustratively.** Cost lines for payments, support and trust & safety are estimates. The purpose is to test structural viability, not to forecast.
- **The competitive gap on guest photos was verified as of the sources cited [R3]** and could close at any time. If Airbnb ships the feature, MVP 1's argument shifts from gap-closing to parity — the loop still functions, but the urgency claim weakens.
- **Frequency of 1–2 trips per user per year is inferred** from Indian leisure-travel patterns, not measured from Airbnb data.

9.3 Limitations — what I would do with more access

The most significant limitation is the absence of **primary research**. With more time I would run 8–12 semi-structured host interviews (focused on the direct-booking decision and willingness to fund advocacy rewards), a 50–100 response guest survey (deposit sensitivity and photo trust), an app-store review teardown across Airbnb, MakeMyTrip and Booking India, and a systematic competitor pricing teardown. Three of the four MVPs rest on behavioural assumptions that qualitative research could de-risk substantially before any experiment is run — and MVP 2's riskiest assumption, whether hosts will part with margin, could very likely be answered in a dozen conversations.

A second limitation: this analysis treats “India” as one market. In practice Goa, Himachal, tier-2 metros and wedding-driven demand behave differently enough that the roadmap would likely need segmenting by cluster rather than sequencing nationally.

SECTION 10

Frameworks, research method & references

10.1 Frameworks applied

FRAMEWORK	WHERE AND WHY IT IS USED
Growth loops (vs. funnels)	Section 5. Chosen over AARRR because a funnel model cannot represent the compounding that makes low-CAC growth possible; loops make the output → input relationship explicit and testable.
Flywheel	Section 5.1. Used to identify <i>where energy must be injected</i> — the analytical step that produces the trust-and-frequency weighting of the whole portfolio.
Product-led growth (applied to internal functions)	Section 4. Deliberately inverted from its usual formulation to answer the brief's actual question about cross-functional enablement.
Lean Startup / MVP & assumption mapping	Section 6. Riskiest-assumption-first ordering, Wizard-of-Oz and painted-door prototypes, and explicit kill criteria.
North Star metric	Section 2.5. Verified trips completed per active traveller per year — chosen to encode all four constraints in one number.
RICE prioritisation	Section 8.1, with an explicit override where the framework's reach term undervalues structural change.
Unit economics & sensitivity analysis	Section 7.3, to test whether the revenue structure survives being wrong about its riskiest input.

10.2 Research method

Research proceeded in three passes. First, a **secondary market pass** establishing India-specific demand structure, competitive commission benchmarks and Airbnb's stated product direction. Second, a **product audit** (Section 3), checking each candidate idea against what Airbnb currently ships — the pass that reclassified two of six ideas. Third, a **build pass**, implementing all four MVPs as a working prototype, which surfaced design questions that were not visible on paper: host right-of-reply placement, the reviews-section-versus-gallery decision, and the need for a control state to make the deposit-waiver argument legible.

Not conducted: primary qualitative or quantitative research with hosts or guests. Section 9.3 sets out what that would have covered and which claims it would have de-risked.

10.3 References

R1	Skift <i>Airbnb on How Indian Travelers Are Staying Local and Spending More</i>	(Sept 2025). Reports 91% domestic share of Airbnb's Indian guests in 2024, up from 79% in 2019; research commissioned by Airbnb and conducted by Oxford Economics.	skift.com/2025/09/09/airbnb-on-how-indian-travelers-are-staying-local-and-spending-more/
R2	MyHomestay <i>OTA Commission Rates India: Cost Guide</i>	. Comparative host commission rates, payout timelines, and estimated annual commission cost to Indian homestays.	myhomestay.ai/blog/ota-commission-guide-india

R3	Skift —	Vrbo Adds Guest Photos to Property Reviews and Galleries	(Oct 2025). Establishes that Booking.com has offered guest photos in reviews since 2016 and that Vrbo launched the feature in October 2025.	skift.com/2025/10/31/vrbo-adds-guest-photos-to-property-review-s-and-galleries/
R4	NichePursuits —	Airbnb Affiliate Program: Why It Ended	. Records the closure of Airbnb's Associates programme on 31 March 2021 with 30 days' notice.	nichepursuits.com/airbnb-affiliate-program/
R5	Annals of Tourism Research (Elsevier, 2025) —	Impact of guest-generated photos: Evidence from Airbnb	. Peer-reviewed evidence on the effect of guest-generated imagery.	sciencedirect.com/science/article/abs/pii/S0160738325000295
R6	RentalScaleUp —	Airbnb Strategy 2025: Key Moves from Q2 Results	. Three strategic pillars; ~94% automated resolution of English-language support enquiries; \$200m investment in Experiences and Services; international expansion elevated to a primary growth pillar.	rentalscaleup.com/airbnb-strategy-2025-q2-results-analysis/
R7	Airbnb Newsroom —	2026 Summer Release	. Services (airport pickups, luggage storage, car rental, grocery), boutique hotels in 20 cities, AI review synthesis, AI support in 11 languages, shared itineraries and the Connections feed.	news.airbnb.com/airbnb-2026-summer-release/
R8	Hostfully —	Airbnb Review Policy Explained: How the Review System Really Works in 2026	. Documents the double-blind two-way review system, the 14-day window, and the host-rates-guest categories.	hostfully.com/blog/airbnb-review-policy/
R9	Airbnb —	Understanding Guest Favorites	, and Axios coverage of the 2023 Winter Release introducing the badge. Establishes existing curation and quality-signalling surfaces.	airbnb.com/resources/hosting-homes/a/understanding-guest-favorites-642
R10	Mordor Intelligence —	India Online Travel Market: Size, Growth & Industry Forecast 2025–2031	. Market sizing and structure context for Section 2.	mordorintelligence.com/industry-reports/online-travel-market-in-india
R11	AirROI —	Top Airbnb Investment Markets in India	. City-level occupancy and revenue structure informing the drive-to and market-concentration observations.	airroi.com/report/world/india

COMPANION ARTEFACT

StayCircle prototype — a single self-contained HTML file implementing all four MVPs (Proof of Stay, Trust Score, Host-Funded Advocacy, Split & Invite) plus a mock experiment dashboard, with a demo-control panel for producing control-versus-variant comparisons. Every screenshot in Section 6 is captured from it. Submitted alongside this document.

The prototype is a concept build created for this coursework. It is not affiliated with, endorsed by, or connected to Airbnb, Inc. All listings, hosts, guests, reviews and metrics within it are fictional.